

Probabilistic Programming Languages

Guillaume Baudart
Christine Tasson

MPRI 2021-2022

MCMC Metropolis-Hastings

Probabilistic Programming Languages

Markov Chain Monte Carlo (MCMC)

Main idea

- Create a Markov chain that converge to the posterior distribution
- Iterate the process until convergence
- Generate samples to approximate the distribution

Pros

- Faster convergence
- Better results for high-dimensional models
- Advanced state-of-the-art optimizations (e.g., HMC, NUTS).

Cons

- Convergences?
- Traps: multimodal, funnel
- Samples correlation

Reminder: Rejection Sampling

coin.ml

```
let coin prob data =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  List.iter (observe prob (bernoulli ~p:z)) data;  
  z  
  
let _ =  
  let d = infer coin [ 1; 1; 0; 0; 0; 0; 0; 0; 0; 0 ] in  
  plot d
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

Reminder: Rejection Sampling


coin.ml

```
let coin prob data =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  List.iter (observe prob (bernoulli ~p:z)) data;  
  z  
  
let _ =  
  let d = infer coin [ 1; 1; 0; 0; 0; 0; 0; 0; 0; 0 ] in  
  plot d
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

10 particles



Reminder: Rejection Sampling


coin.ml

```
let coin prob data =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  List.iter (observe prob (bernoulli ~p:z)) data;  
  z  
  
let _ =  
  let d = infer coin [ 1; 1; 0; 0; 0; 0; 0; 0; 0; 0 ] in  
  plot d
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

100 particles



Reminder: Rejection Sampling


coin.ml

```
let coin prob data =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  List.iter (observe prob (bernoulli ~p:z)) data;  
  z  
  
let _ =  
  let d = infer coin [ 1; 1; 0; 0; 0; 0; 0; 0; 0; 0 ] in  
  plot d
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

1000 particles



Reminder: Rejection Sampling


coin.ml

```
let coin prob data =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  List.iter (observe prob (bernoulli ~p:z)) data;  
  z  
  
let _ =  
  let d = infer coin [ 1; 1; 0; 0; 0; 0; 0; 0; 0; 0 ] in  
  plot d
```

Executing the model generates one sample

- **sample**: draw from a distribution
- **assume/observe**: hard conditioning, reject invalid samples
- Terminates with *n* valid samples

1000 particles



Weighted Rejection Sampling

Adapt rejection sampling to soft conditioning

- Execute the sampler to get a pair (v_i, w_i)
- Suppose w_{\max} is known
- Accept the sample with probability w_i/w_{\max} or retry

Weighted Rejection Sampling

Adapt rejection sampling to soft conditioning

- Execute the sampler to get a pair (v_i, w_i)
- Suppose w_{\max} is known
- Accept the sample with probability w_i/w_{\max} or retry

But w_{\max} is not known...

Execution Trace

Consider a program execution with

- $X = x_0, \dots, x_n$: set of random variables sampled at step i : the trace
- $Y = y_0, \dots, y_m$: set of random variables observed at step i .

Remarks

- Sets X and Y depend on the execution path
- We can only control X

Execution Trace

Consider a program execution with

- $X = x_0, \dots, x_n$: set of random variables sampled at step i : the trace
- $Y = y_0, \dots, y_m$: set of random variables observed at step i .

Remarks

- Sets X and Y depend on the execution path
- We can only control X

```
let bimodal y =  
  let z = sample (bernoulli ~p:0.5) in  
  let mu =  
    if z then sample (gaussian ~mu:-1. ~sigma:1.)  
    else sample (gaussian ~mu:1. ~sigma:1.)  
  in  
  let () = observe (gaussian ~mu ~sigma:1.) y in  
  z
```

Multi-Sites Metropolis Hastings


Markov chain on execution traces

- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}

Multi-Sites Metropolis Hastings

Markov chain on execution traces


- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}



Multi-Sites Metropolis Hastings

Markov chain on execution traces

- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}



outputs


v_0

v_1

Multi-Sites Metropolis Hastings

Markov chain on execution traces

- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}



outputs

v_0


v_1

v_1

Multi-Sites Metropolis Hastings

Markov chain on execution traces

- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}



outputs

v_0

v_1

v_1

v_3

...

General Metropolis Hastings

More generally

- $X = x_0, \dots, x_n$: set of random variables sampled at step i : the trace
- $Y = y_0, \dots, y_m$: set of random variables observed at step i .
- Propose a new trace from a proposal distribution $q(X_i | X_{i-1})$
- Accept the trace with probability α , where

$$\alpha = \min \left(1, \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} \right)$$

- Otherwise return the previous trace X_{i-1}

Multi-Sites Metropolis Hastings: Acceptation

- Draw proposal from priors: $q(X_i | X_{i-1}) = p(X_i)$
- Resample all variable in X_i at each step

$$\begin{aligned} \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} &= \frac{p(Y_i | X_i) p(X_i)}{p(Y_{i-1} | X_{i-1}) p(X_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} \\ &= \frac{p(Y_i | X_i) p(X_i)}{p(Y_{i-1} | X_{i-1}) p(X_{i-1})} \frac{p(X_{i-1})}{p(X_i)} \\ &= \frac{p(Y_i | X_i)}{p(Y_{i-1} | X_{i-1})} \\ &= \frac{w_i}{w_{i-1}} \end{aligned}$$

Multi-Sites Metropolis Hastings: Acceptation

- Draw proposal from priors: $q(X_i | X_{i-1}) = p(X_i)$
- Resample all variable in X_i at each step

$$\begin{aligned} \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} &= \frac{p(Y_i | X_i) p(X_i)}{p(Y_{i-1} | X_{i-1}) p(X_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} \\ &= \frac{p(Y_i | X_i) p(X_i)}{p(Y_{i-1} | X_{i-1}) p(X_{i-1})} \frac{p(X_{i-1})}{p(X_i)} \\ &= \frac{p(Y_i | X_i)}{p(Y_{i-1} | X_{i-1})} \\ &= \frac{w_i}{w_{i-1}} \end{aligned}$$

Markov chain on execution traces

- Execute the sampler to get a trace and associated score (X_i, w_i)
- If $w_i \geq w_{i-1}$ accept the trace (and the associated output)
- Else accept the trace with probability w_i/w_{i-1}
- Otherwise return the previous trace X_{i-1}

Single-Site Metropolis Hastings


Reuse most of the previous trace (i.e., sampled values)

- Choose one random variable to resample to obtain a new execution
- Accept the trace with probability α
- Otherwise use the previous trace

Single-Site Metropolis Hastings

Reuse most of the previous trace (i.e., sampled values)


- Choose one random variable to resample to obtain a new execution
- Accept the trace with probability α
- Otherwise use the previous trace



Single-Site Metropolis Hastings

Reuse most of the previous trace (i.e., sampled values)


- Choose one random variable to resample to obtain a new execution
- Accept the trace with probability α
- Otherwise use the previous trace



Single-Site Metropolis Hastings

Reuse most of the previous trace (i.e., sampled values)


- Choose one random variable to resample to obtain a new execution
- Accept the trace with probability α
- Otherwise use the previous trace



Single-Site Metropolis Hastings

Reuse most of the previous trace (i.e., sampled values)

- Choose one random variable to resample to obtain a new execution
- Accept the trace with probability α
- Otherwise use the previous trace



Single Site Metropolis Hastings: Acceptation

Track the likelihood of all random variable during execution

- $x = \text{sample } d \rightarrow w(x) = (x, \text{pdf } d \ x)$
- $\text{observe } d \ y \rightarrow w(y) = (y, \text{pdf } d \ y)$

Single Site Metropolis Hastings: Acceptation

Track the likelihood of all random variable during execution

- `x = sample d` $\rightarrow w(x) = (x, \text{pdf } d \ x)$
- `observe d y` $\rightarrow w(y) = (y, \text{pdf } d \ y)$

```
let bimodal y =  
  let z = sample (bernoulli ~p:0.5) in  
  let mu =  
    if z then sample (gaussian ~mu:-1. ~sigma:1.)  
    else sample (gaussian ~mu:1. ~sigma:1.)  
  in  
  let () = observe (gaussian ~mu ~sigma:1.) y in  
  z
```

Single Site Metropolis Hastings: Acceptation

Track the likelihood of all random variable during execution

- `x = sample d` $\rightarrow w(x) = (x, \text{pdf } d \ x)$
- `observe d y` $\rightarrow w(y) = (y, \text{pdf } d \ y)$

```
let bimodal y =  
  let z = sample (bernoulli ~p:0.5) in  
  let mu =  
    if z then sample (gaussian ~mu:-1. ~sigma:1.)  
    else sample (gaussian ~mu:1. ~sigma:1.)  
  in  
  let () = observe (gaussian ~mu ~sigma:1.) y in  
  z
```

```
w(z)   = (false, 0.5)  
w(mu)  = (1.2, 0.40)  
w(y)   = (2.0, 0.27)
```

Single Site Metropolis Hastings: Acceptation

Track the likelihood of all random variable during execution

- `x = sample d` $\rightarrow w(x) = (x, \text{pdf } d \ x)$
- `observe d y` $\rightarrow w(y) = (y, \text{pdf } d \ y)$ score (as in importance sampling)

```
let bimodal y =  
  let z = sample (bernoulli ~p:0.5) in  
  let mu =  
    if z then sample (gaussian ~mu:-1. ~sigma:1.)  
    else sample (gaussian ~mu:1. ~sigma:1.)  
  in  
  let () = observe (gaussian ~mu ~sigma:1.) y in  
  z
```

```
w(z)    = (false, 0.5)  
w(mu)   = (1.2, 0.40)  
w(y)    = (2.0, 0.27)
```

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$\frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} = \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})}$$

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$\frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} = \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})}$$

$$p(X, Y) = \prod_{x \in X} w(x) \prod_{y \in Y} w(y) \quad (\text{not necessarily independent})$$

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$\frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} = \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})}$$

$$p(X, Y) = \prod_{x \in X} w(x) \prod_{y \in Y} w(y) \quad (\text{not necessarily independent})$$

$$q(x_0 | X) = 1/|X| \quad 1/(\text{n choices})$$

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$\frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} = \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})}$$

$$p(X, Y) = \prod_{x \in X} w(x) \prod_{y \in Y} w(y) \quad \text{(not necessarily independent)}$$

$$q(x_0 | X) = 1/|X| \quad \text{1/(n choices)}$$

$$q(X | X', x_0) = \prod_{x \in X^{\text{sample}}} w(x) \quad \text{Resampled from the prior}$$

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$\begin{aligned}
 \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} &= \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})} \\
 &= \frac{q(x_0 | X_i)}{q(x_0 | X_{i-1})} \frac{p(X_i, Y_i)}{q(X_i | X_{i-1}, x_0)} \frac{q(X_{i-1} | X_i, x_0)}{p(X_{i-1}, Y_{i-1})} \\
 &= \frac{|X_{i-1}|}{|X_i|} \frac{\prod_{x \in X_i} w(x) \prod_{y \in Y_i} w(y)}{\prod_{x \in X_i^{\text{sample}}} w(x)} \frac{\prod_{x \in X_{i-1}^{\text{sample}}} w(x)}{\prod_{x \in X_{i-1}} w(x) \prod_{y \in Y_{i-1}} w(y)} \\
 &= \frac{|X_{i-1}|}{|X_i|} \frac{\prod_{x \in X_i^{\text{reuse}}} w(x)}{\prod_{x \in X_{i-1}^{\text{reuse}}} w(x)} \frac{\prod_{y \in Y_i} w(y)}{\prod_{y \in Y_{i-1}} w(y)}
 \end{aligned}$$

$$p(X, Y) = \prod_{x \in X} w(x) \prod_{y \in Y} w(y)$$

$$q(x_0 | X) = |X|$$

$$q(X | X', x_0) = \prod_{x \in X^{\text{sample}}} w(x)$$

Single Site Metropolis Hastings: Acceptation

- Pick one variable x_0 at random in the trace
- Resample only this variable and all its dependencies $X = X^{\text{sample}} \cup X^{\text{reuse}}$

$$p(X, Y) = \prod_{x \in X} w(x) \prod_{y \in Y} w(y)$$

$$q(x_0 | X) = |X|$$

$$q(X | X', x_0) = \prod_{x \in X^{\text{sample}}} w(x)$$

$$\begin{aligned}
 \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i)}{q(X_i | X_{i-1})} &= \frac{p(X_i, Y_i)}{p(X_{i-1}, Y_{i-1})} \frac{q(X_{i-1} | X_i, x_0) q(x_0 | X_i)}{q(X_i | X_{i-1}, x_0) q(x_0 | X_{i-1})} \\
 &= \frac{q(x_0 | X_i)}{q(x_0 | X_{i-1})} \frac{p(X_i, Y_i)}{q(X_i | X_{i-1}, x_0)} \frac{q(X_{i-1} | X_i, x_0)}{p(X_{i-1}, Y_{i-1})} \\
 &= \frac{|X_{i-1}|}{|X_i|} \frac{\prod_{x \in X_i} w(x) \prod_{y \in Y_i} w(y)}{\prod_{x \in X_i^{\text{sample}}} w(x)} \frac{\prod_{x \in X_{i-1}^{\text{sample}}} w(x)}{\prod_{x \in X_{i-1}} w(x) \prod_{y \in Y_{i-1}} w(y)} \\
 &= \frac{|X_{i-1}|}{|X_i|} \frac{\prod_{x \in X_i^{\text{reuse}}} w(x)}{\prod_{x \in X_{i-1}^{\text{reuse}}} w(x)} \frac{\prod_{y \in Y_i} w(y)}{\prod_{y \in Y_{i-1}} w(y)} \\
 &\quad \begin{array}{ccc} \downarrow & \downarrow & \downarrow \\ \text{choice } x_0 & \text{reused} & \text{scores} \end{array}
 \end{aligned}$$

Implementation

Implementation

Option 1

- CPS model
- Pick x_0 at random
- Resample all variables after x_0

Implementation

Option 1

- CPS model
- Pick x_0 at random
- Resample all variables after x_0

Option 2

- CPS model
- Pick x_0 at random
- Memoize values of all random variable
- Resample only when needed

Implementation

Option 1

- CPS model
- Pick x_0 at random
- Resample all variables after x_0

Option 2

- CPS model
- Pick x_0 at random
- Memoize values of all random variable
- Resample only when needed

Warning: Memoization naming scheme
Different variable may have the same name...

Limitations

Limitations

Convergence

- Theoretical conditions are complex
- Must be checked experimentally
- Diagnostic tools: trace plot, R-hat (multi-chains)
- Solution: warmup, change initial conditions, reparameterization, ...

Limitations

Convergence

- Theoretical conditions are complex
- Must be checked experimentally
- Diagnostic tools: trace plot, R-hat (multi-chains)
- Solution: warmup, change initial conditions, reparameterization, ...

Sample correlation

- Next sample depends on the previous one
- Diagnostic tools ESS (effective sample size)
- Solution: thinning, ...

Limitations

Convergence

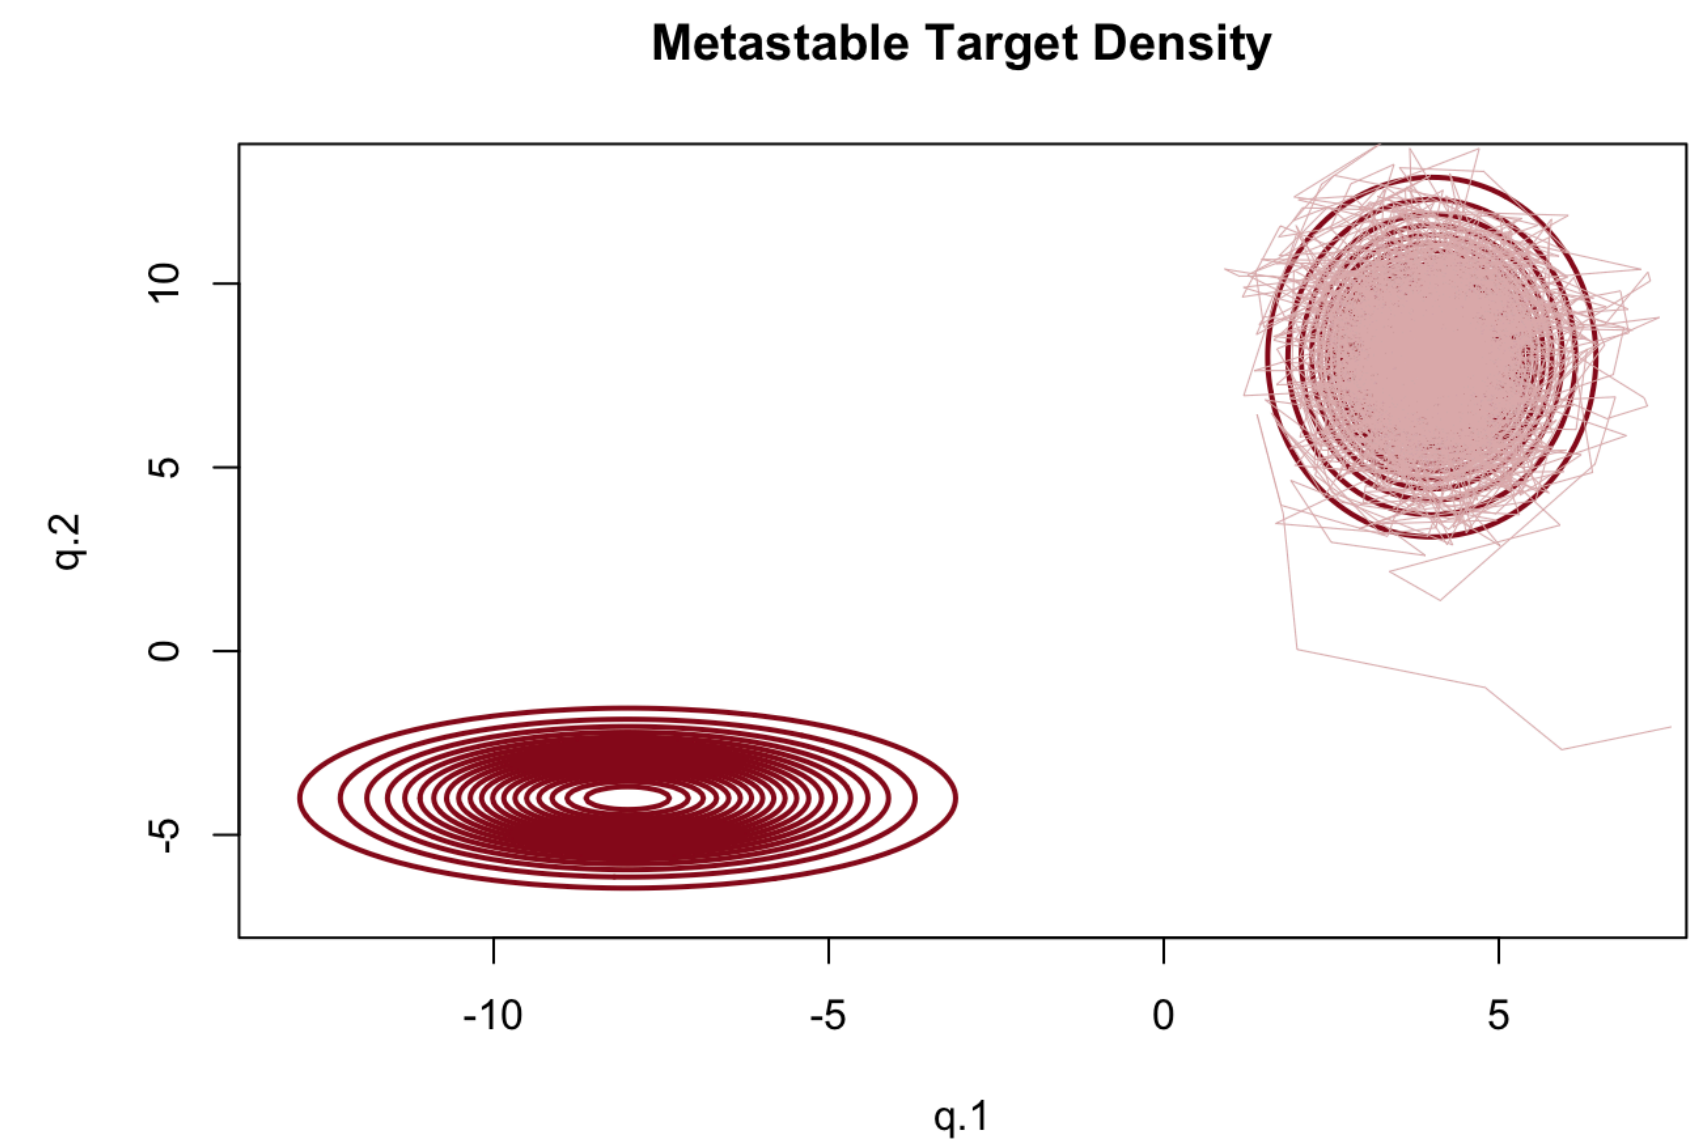
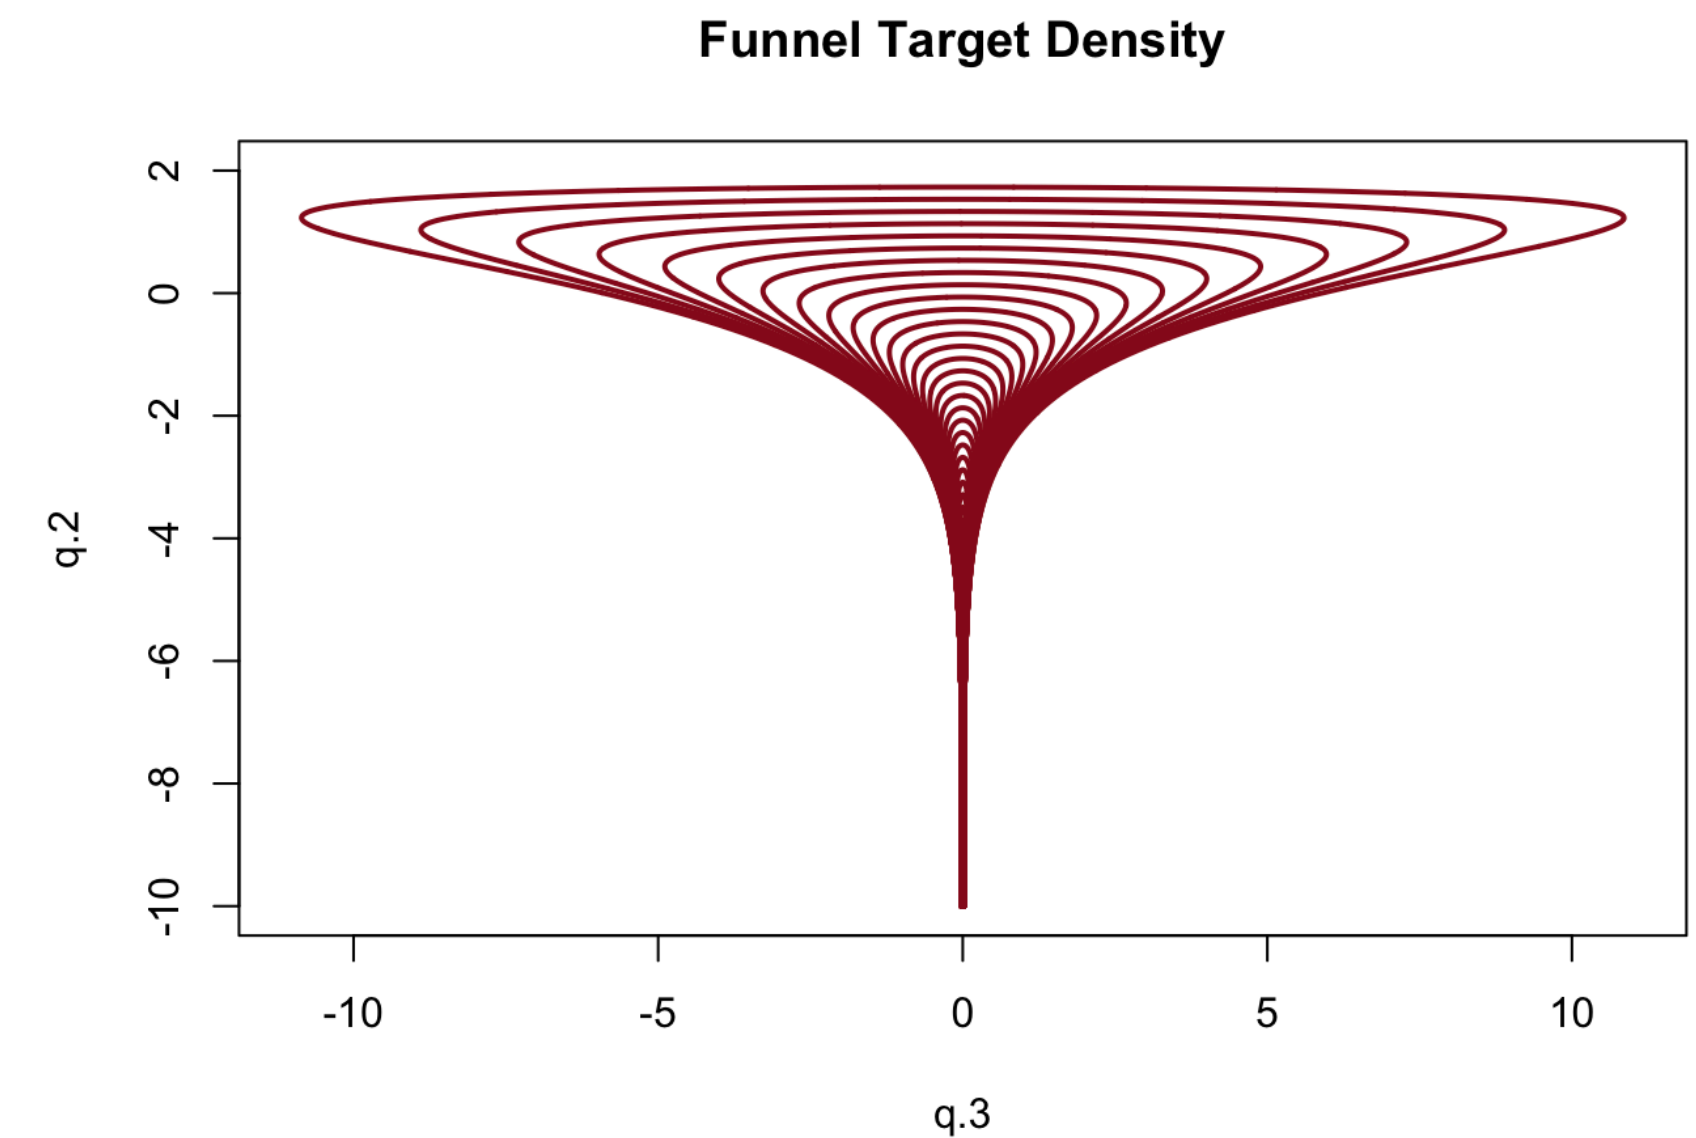
- Theoretical conditions are complex
- Must be checked experimentally
- Diagnostic tools: trace plot, R-hat (multi-chains)
- Solution: warmup, change initial conditions, reparameterization, ...

Sample correlation

- Next sample depends on the previous one
- Diagnostic tools ESS (effective sample size)
- Solution: thinning, ...

Pathological models

- Multimodal distribution
- Neal's Funnel



Advanced Inference

Probabilistic Programming Languages

Hamiltonian Monte-Carlo (HMC)

Hamiltonian Monte-Carlo (HMC)

Analogy: Particle in an energy field

- Program define a density of the form $p(X) \propto \exp(-U(X))$
- On continuous spaces U can be interpreted as an energy
- Low energy wells correspond to high probability regions
- HMC simulate the trajectory of a particle in this energy field

Hamiltonian Monte-Carlo (HMC)

Analogy: Particle in an energy field

- Program define a density of the form $p(X) \propto \exp(-U(X))$
- On continuous spaces U can be interpreted as an energy
- Low energy wells correspond to high probability regions
- HMC simulate the trajectory of a particle in this energy field

Hamiltonian dynamic

- M : Mass matrix
- P : Momentum
- $K(P) = \frac{1}{2}P^T M P$

Hamiltonian Monte-Carlo (HMC)

Analogy: Particle in an energy field

- Program define a density of the form $p(X) \propto \exp(-U(X))$
- On continuous spaces U can be interpreted as an energy
- Low energy wells correspond to high probability regions
- HMC simulate the trajectory of a particle in this energy field

Hamiltonian dynamic

- M : Mass matrix
- P : Momentum
- $K(P) = \frac{1}{2}P^T M P$

The diagram illustrates the components of the Hamiltonian. The equation $H(X, P) = K(P) + U(X)$ is centered. An upward arrow from $H(X, P)$ points to the word "hamiltonian". Another upward arrow from $U(X)$ points to the words "potential energy". A downward arrow from $K(P)$ points to the words "kinetic energy".

$$\begin{array}{ccc} \text{hamiltonian} & & \text{potential energy} \\ \uparrow & & \uparrow \\ H(X, P) = K(P) + U(X) \\ \downarrow & & \\ \text{kinetic energy} & & \end{array}$$

Hamiltonian Monte-Carlo (HMC)

The diagram illustrates the components of the Hamiltonian function. The central equation is $H(X, P) = K(P) + U(X)$. An upward arrow from $H(X, P)$ points to the label "hamiltonian". A downward arrow from $K(P)$ points to the label "kinetic energy". An upward arrow from $U(X)$ points to the label "potential energy".

hamiltonian

potential energy

$$H(X, P) = K(P) + U(X)$$

kinetic energy

Hamiltonian Monte-Carlo (HMC)

Energy conservation

$$\frac{dH}{dt} = (\nabla_P H)^T \frac{dP}{dt} + (\nabla_X H)^T \frac{dX}{dt} = 0$$

A diagram illustrating the components of the Hamiltonian. The central equation is $H(X, P) = K(P) + U(X)$. An upward arrow from $K(P)$ points to the word "kinetic energy". An upward arrow from $U(X)$ points to the word "potential energy". A downward arrow from the entire equation $H(X, P) = K(P) + U(X)$ points to the word "hamiltonian".

$$H(X, P) = K(P) + U(X)$$

kinetic energy

hamiltonian

potential energy

Hamiltonian Monte-Carlo (HMC)

Energy conservation

$$\frac{dH}{dt} = (\nabla_P H)^T \frac{dP}{dt} + (\nabla_X H)^T \frac{dX}{dt} = 0$$

Hamiltonian dynamics

$$\begin{cases} \frac{dX}{dt} = \nabla_P H(X, P) = M^{-1}P \\ \frac{dP}{dt} = -\nabla_X H(X, P) = -\nabla_X U(X) \end{cases}$$

The diagram illustrates the decomposition of the Hamiltonian $H(X, P)$ into its components. The equation $H(X, P) = K(P) + U(X)$ is centered. An upward arrow from $H(X, P)$ points to the label "hamiltonian". Another upward arrow from $U(X)$ points to the label "potential energy". A downward arrow from $K(P)$ points to the label "kinetic energy".

$$\begin{array}{ccc} \text{hamiltonian} & & \text{potential energy} \\ \uparrow & & \uparrow \\ H(X, P) = K(P) + U(X) \\ \downarrow & & \\ \text{kinetic energy} & & \end{array}$$

Hamiltonian Monte-Carlo (HMC)

Hamiltonian Monte-Carlo (HMC)

Generate samples from $p(X, P) \propto \exp(-H(X, P))$

- At each iteration
- Sample an initial momentum $P_i(0) \sim \mathcal{N}(0, M)$
- Solve the Hamiltonian dynamic (discretized)
- Perform a Metropolis Hastings update with probability α

$$\alpha = \min \left(1, \frac{\exp(-H(X_i, P_i))}{\exp(-H(X_{i-1}, P_{i-1}))} \right)$$

Hamiltonian Monte-Carlo (HMC)

Generate samples from $p(X, P) \propto \exp(-H(X, P))$

- At each iteration
- Sample an initial momentum $P_i(0) \sim \mathcal{N}(0, M)$
- Solve the Hamiltonian dynamic (discretized)
- Perform a Metropolis Hastings update with probability α

$$\alpha = \min \left(1, \frac{\exp(-H(X_i, P_i))}{\exp(-H(X_{i-1}, P_{i-1}))} \right)$$

If the hamiltonian is preserved: accept with probability 1.

Hamiltonian Monte-Carlo (HMC)

Generate samples from $p(X, P) \propto \exp(-H(X, P))$


- At each iteration
- Sample an initial momentum $P_i(0) \sim \mathcal{N}(0, M)$
- Solve the Hamiltonian dynamic (discretized)
- Perform a Metropolis Hastings update with probability α

$$\alpha = \min \left(1, \frac{\exp(-H(X_i, P_i))}{\exp(-H(X_{i-1}, P_{i-1}))} \right)$$

If the hamiltonian is preserved: accept with probability 1.


Parameter P can be marginalized

Stochastic Variational Inference (SVI)



Stochastic Variational Inference (SVI)

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int_z p(x|z)p(z)dz}$$




Stochastic Variational Inference (SVI)

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int_z p(x|z)p(z)dz}$$

Variational family

- Parameterized by a parameter ϕ
- Find the closest member to the posterior $q_{\phi^*}(z)$
- Optimization problem



Stochastic Variational Inference (SVI)

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int_z p(x|z)p(z)dz}$$


Variational family

- Parameterized by a parameter ϕ
- Find the closest member to the posterior $q_{\phi^*}(z)$
- Optimization problem

Metrics: Kullback-Leibler divergence

$$KL(q(x) || p(x)) = - \int q(x) \log \frac{p(x)}{q(x)} dx$$

- $KL(q || p) \geq 0$
- $KL(q || p) \neq KL(p || q)$



Stochastic Variational Inference (SVI)

$$\begin{aligned} KL(q_\phi(z) \parallel p(z|x)) &= - \int q_\phi(z) \log \frac{p(z|x)}{q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{p(x)q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \int q_\phi(z) \log p(x) dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \log p(x) \end{aligned}$$

Stochastic Variational Inference (SVI)


$$\begin{aligned} KL(q_\phi(z) || p(z|x)) &= - \int q_\phi(z) \log \frac{p(z|x)}{q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{p(x)q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \int q_\phi(z) \log p(x) dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \log p(x) \end{aligned}$$

$$\log p(x) = KL(q_\phi(z) || p(x, z)) + \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz$$

Stochastic Variational Inference (SVI)

$$\begin{aligned} KL(q_\phi(z) || p(z|x)) &= - \int q_\phi(z) \log \frac{p(z|x)}{q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{p(x)q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \int q_\phi(z) \log p(x) dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \log p(x) \end{aligned}$$

$$\log p(x) = KL(q_\phi(z) || p(x, z)) + \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz$$




constant


Stochastic Variational Inference (SVI)

$$\begin{aligned} KL(q_\phi(z) || p(z|x)) &= - \int q_\phi(z) \log \frac{p(z|x)}{q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{p(x)q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \int q_\phi(z) \log p(x) dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \log p(x) \end{aligned}$$

$$\log p(x) = KL(q_\phi(z) || p(x, z)) + \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz$$



constant



minimize

Stochastic Variational Inference (SVI)

$$\begin{aligned} KL(q_\phi(z) || p(z|x)) &= - \int q_\phi(z) \log \frac{p(z|x)}{q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{p(x)q_\phi(z)} dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \int q_\phi(z) \log p(x) dz \\ &= - \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz + \log p(x) \end{aligned}$$

$$\log p(x) = KL(q_\phi(z) || p(x, z)) + \int q_\phi(z) \log \frac{p(x, z)}{q_\phi(z)} dz$$

constant

minimize

maximize ELBO \mathcal{L}

Stochastic Variational Inference (SVI)

$$\begin{aligned}\mathcal{L} &= \int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz \\ &= \int q_{\phi}(z) \log \frac{p(x|z)p(z)}{q_{\phi}(z)} dz \\ &= \int q_{\phi}(z) \log p(x|z) dz + \int q_{\phi}(z) \log \frac{p(z)}{q_{\phi}(z)} dz\end{aligned}$$

$$\log p(x) = \underbrace{KL(q_{\phi}(z) || p(x, z))}_{\text{constant}} + \underbrace{\int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz}_{\text{minimize}} + \underbrace{\int q_{\phi}(z) \log \frac{p(z)}{q_{\phi}(z)} dz}_{\text{maximize ELBO } \mathcal{L}}$$

Stochastic Variational Inference (SVI)

$$\begin{aligned}\mathcal{L} &= \int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz \\ &= \int q_{\phi}(z) \log \frac{p(x|z)p(z)}{q_{\phi}(z)} dz \\ &= \underbrace{\int q_{\phi}(z) \log p(x|z) dz}_{\substack{\text{maximize likelihood} \\ \mathbb{E}_q \log p(x|z)}} + \int q_{\phi}(z) \log \frac{p(z)}{q_{\phi}(z)} dz\end{aligned}$$

$$\log p(x) = \underbrace{KL(q_{\phi}(z) || p(x, z))}_{\text{constant}} + \underbrace{\int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz}_{\text{minimize}} + \underbrace{\int q_{\phi}(z) \log \frac{p(z)}{q_{\phi}(z)} dz}_{\text{maximize ELBO } \mathcal{L}}$$

Stochastic Variational Inference (SVI)

$$\begin{aligned}\mathcal{L} &= \int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz \\ &= \int q_{\phi}(z) \log \frac{p(x|z)p(z)}{q_{\phi}(z)} dz \\ &= \underbrace{\int q_{\phi}(z) \log p(x|z) dz}_{\substack{\text{maximize likelihood} \\ \mathbb{E}_q \log p(x|z)}} + \underbrace{\int q_{\phi}(z) \log \frac{p(z)}{q_{\phi}(z)} dz}_{\substack{\text{minimize divergence} \\ -KL(q(z) || p(z))}}\end{aligned}$$

$$\log p(x) = \underbrace{KL(q_{\phi}(z) || p(x, z))}_{\text{constant}} + \underbrace{\int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz}_{\text{minimize}} + \underbrace{\int q_{\phi}(z) \log \frac{p(x, z)}{q_{\phi}(z)} dz}_{\text{maximize ELBO } \mathcal{L}}$$

Variational Family

Variational Family

Black-box variational inference

- Mean-field approximation $q_{\phi}(z) = \prod_{i=1}^n \mathcal{N}(z_i | \mu_i, \sigma_i)$ where $\phi = \{\mu_i, \sigma_i\}_{i \in [1, n]}$
- Full-rank approximation $q_{\phi}(z) = \mathcal{N}(z | \mu, \Sigma)$ where $\phi = (\mu, \Sigma)$
- Pyro autoguides

Variational Family

Black-box variational inference

- Mean-field approximation $q_{\phi}(z) = \prod_{i=1}^n \mathcal{N}(z_i | \mu_i, \sigma_i)$ where $\phi = \{\mu_i, \sigma_i\}_{i \in [1, n]}$
- Full-rank approximation $q_{\phi}(z) = \mathcal{N}(z | \mu, \Sigma)$ where $\phi = (\mu, \Sigma)$
- Pyro autoguides

Program your own guide

- Pyro (first versions)
- Must sample the same variables in the guide and the model
- Static analysis?

```
def model():  
    pyro.sample("z_1", ... )  
  
def guide():  
    pyro.sample("z_1", ... )
```

References

An Introduction to Probabilistic Programming

Jan-Willem van de Meent, Brooks Paige, Hongseok Yang, Frank Wood

<https://arxiv.org/abs/1809.10756>

Lightweight Implementations of Probabilistic Programming Languages Via Transformational Compilation

David Wingate Andreas Stuhlmüller Noah D. Goodman

AISTATS 2011

Markov Chain Monte Carlo in Practice

Michael Bettancourt

https://betanalpha.github.io/assets/case_studies/markov_chain_monte_carlo.html

Variational Inference: A Review for Statisticians

David M. Blei, Alp Kucukelbir, Jon D. McAuliffe

<https://arxiv.org/abs/1601.00670>

Compiling Stan to Generative Probabilistic Languages and Extension to Deep Probabilistic Programming

Guillaume Baudart, Javier Burroni, Martin Hirzel, Louis Mandel, Avraham Shinnar

PLDI 2021